

---

# SongRNN: Music Generation Through a Character Level LSTM

---

**Brandon Szeto**

Jacobs School of Engineering  
University of California, San Diego  
San Diego, CA 92122  
bszeto@ucsd.edu

**Darren Yu**

Jacobs School of Engineering  
University of California, San Diego  
San Diego, CA 92122  
dmyu@ucsd.edu

**Nathaniel Thomas**

Jacobs School of Engineering  
University of California, San Diego  
San Diego, CA 92122  
nathomas@ucsd.edu

## Abstract

With a set of musical samples in ABC notation, we can generate music with a character-level LSTM. We can accomplish this task by feeding our LSTM musical characters in a sequence, and the network would slowly learn to predict the upcoming sequence of notes. Once trained, our model could generate music samples based on a short sequence of prompted notes. The style of the music generated by our final model has a rich, mature, and elaborate feel. By modifying the network depth, layer size, and dropout rate, we could tune our model to generate a range of musical samples. After hyperparameter tuning, we found that our LSTM created more elegant musical samples than a traditional recurrent neural network. The ability of the LSTM to drop information allows it to resemble the original test samples better than a standard RNN.

## Introduction

From classical era music to modern hip-hop, music has been ingrained in human culture for thousands of years. The idea of creating music is assumed to be a task only accomplishable by humans. But with a large sample of composed music, a recurrent network model, and some computational power, we can generate music without the interventions of humans. A long short-term memory network (LSTM) is an type of RNN that uses gates to control the flow of information in and out of the network. LSTMs are great at predicting sequences of information based on input information. With this in mind, our goal is to generate music from a short prompt of musical characters. We can accomplish this task with a character-level LSTM.

## Related Works

### SongRNN

Briot et al. [2019]

## Methods

In the Methods section, you should describe the implementation and architectural details of your system - in particular, this addresses how you approached the problem and the design of your solution. For those who believe in reproducible science, this should be fine-grained enough such that somebody could implement your model/algorithm and reproduce the results you claim to achieve.

### Training network using Teacher forcing

Describe the LSTM architecture, the training procedure, the loss criterion, the optimizer you used.

### Song Generation

Describe the approach you took to generate a song by priming the network with an initial sequence given to it. You should also discuss the details of incorporating Temperature (T) in generating the output.

### Hyper-parameter Tuning

Describe the RNN architecture you used in 5.a and briefly describe the approach you took in tuning your hyperparameters.

### Feature Evaluation

Describe why it is important to do feature evaluation. Describe the approach you took in generating the Heatmap of the activation's of each neurons for each of the characters.

## Results

In the Results section, you should demonstrate your baseline model's performance, performances for each of the parts of Q3, Q4, Q5 and Q6 (training, generation, hyperparameter tuning, feature evaluation). You are expected to report the following things in your submission:

- a Report results for your baseline model, a 1-hidden layer LSTM with 150 neurons. (2 pt)
- b Report the three generated music samples with the above-mentioned three different Temperature (T) settings. Submit their .txt and .midi files as a part of your gradescope submission. (3 pt)
- c Report loss plots for hyperparameter tuning experiments namely: RNNs vs LSTMs, changing number of neurons in the hidden layer and varying dropouts. (6 pt)

- d Report heatmaps with appropriate scales, for three different neurons from your best-trained model. (3 pt)

## Discussion

Please discuss the following important points as well: How did the qualitative performance and loss of RNN vs LSTM models differ? Discuss the performance differences with respect to changes in the number of neurons and dropout. Also, draw insights from the heatmap of each of your neurons.

## Contributions

**Brandon Szeto:** Data evaluation, loss graphs, diagrams, write up (related work, methods, and discussion).

**Darren Yu:** LSTM/RNN hyperparameter tuning, feature evaluation, and write up (abstract, introduction, and results),

**Nathaniel Thomas:** LSTM, SongRNN, model architecture, model training, and music generation.

## References

Jean-Pierre Briot, Gaëtan Hadjeres, and François-David Pachet. Deep learning techniques for music generation – a survey, 2019.